

# Posterior Predictive Checks: Reaction Time Example

brms Workshop

## Posterior Predictive Checks for RT Data

After fitting your model, validate that it generates data similar to what you observed.

### Why Posterior Predictive Checks Matter

Posterior predictive checks answer: “If I were to generate new data from my fitted model, would it look like my actual data?” This validates that your model has captured the essential structure of your data.

### Setup

```
library(brms)
library(tidyverse)
library(bayesplot)

# Set seed for reproducibility
set.seed(42)
```

### Load or Fit Model

```
# Create example RT data (same as in prior checks for consistency)
n_subj <- 20
n_trials <- 50
n_items <- 30
```

```

rt_data <- expand.grid(
  trial = 1:n_trials,
  subject = 1:n_subj,
  item = 1:n_items
) %>%
  filter(row_number() <= n_subj * n_trials * 3) %>%
  mutate(
    condition = rep(c("A", "B"), length.out = n()),
    log_rt = rnorm(n(), mean = 6, sd = 0.3) +
      (condition == "B") * 0.15 +
      rnorm(n(), mean = 0, sd = 0.1),
    rt = exp(log_rt)
  )

# Define priors
rt_priors <- c(
  prior(normal(6, 1.5), class = Intercept),
  prior(normal(0, 0.5), class = b),
  prior(exponential(1), class = sigma),
  prior(exponential(1), class = sd),
  prior(lkj(2), class = cor)
)

# Check if model exists, otherwise fit it
model_file <- "fits/fit_rt.rds"
if (file.exists(model_file)) {
  cat("Loading saved model from:", model_file, "\n")
  fit_rt <- readRDS(model_file)
} else {
  cat("Fitting model (this may take a while)... \n")
  cat("Note: Model fitting requires significant computational resources.\n")
  cat("Consider fitting the model separately and saving to fits/fit_rt.rds\n\n")

  # Fit with reduced complexity for demonstration
  fit_rt <- brm(
    log_rt ~ condition + (1 + condition | subject) + (1 | item),
    data = rt_data,
    family = gaussian(),
    prior = rt_priors,
    chains = 2, # Reduced for faster fitting
    iter = 1000,
    cores = 2,
  )
}

```

```

    backend = "rstan",
    refresh = 0
)
# Save for future use
dir.create("fits", showWarnings = FALSE, recursive = TRUE)
saveRDS(fit_rt, model_file)
cat("Model saved to:", model_file, "\n")
}

```

Loading saved model from: fits/fit\_rt.rds

```

# Display model summary
cat("\n==== Model Summary ===\n")

```

==== Model Summary ===

```
print(summary(fit_rt))
```

```

Family: gaussian
Links: mu = identity
Formula: log_rt ~ condition + (1 + condition | subject) + (1 | item)
Data: rt_data (Number of observations: 3000)
Draws: 2 chains, each with iter = 1000; warmup = 500; thin = 1;
      total post-warmup draws = 1000

```

Multilevel Hyperparameters:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
~item (Number of levels: 3)							
sd(Intercept)	0.02	0.02	0.00	0.07	1.01	392	323
~subject (Number of levels: 20)							
sd(Intercept)	0.02	0.01	0.00	0.04	1.00	432	
sd(conditionB)	0.02	0.01	0.00	0.06	1.00	343	
cor(Intercept,conditionB)	-0.07	0.46	-0.82	0.83	1.01	513	
					Tail_ESS		
sd(Intercept)		528					
sd(conditionB)		461					
cor(Intercept,conditionB)		591					

Regression Coefficients:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	5.99	0.01	5.96	6.02	1.00	700	436
conditionB	0.16	0.01	0.13	0.19	1.00	1339	699

Further Distributional Parameters:

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	0.32	0.00	0.31	0.33	1.01	2149	698


Draws were sampled using sampling(NUTS). For each parameter, Bulk\_ESS and Tail\_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

## Basic Posterior Predictive Checks

### Visual Checks

```
# Default: density overlay of observed vs. simulated data
pp_check(fit_rt, ndraws = 100) +
  ggtitle("Density overlay: Observed vs. Posterior predictions")
```

Density overlay: Observed vs. Posterior predictions



### Interpretation:

- Blue line (observed data) should be among the dark lines (posterior predictions)
- If blue line is far from the bundle → model missed something important
- Small discrepancies are normal; large ones suggest model misspecification

### Check Specific Statistics

```
# Did we get the mean right?  
p1 <- pp_check(fit_rt, type = "stat", stat = "mean") +  
  ggtitle("Mean: Observed vs. Predicted")  
  
# Did we get the spread right?  
p2 <- pp_check(fit_rt, type = "stat", stat = "sd") +
```

```


ggttitle("SD: Observed vs. Predicted")

# Extreme values?
p3 <- pp_check(fit_rt, type = "stat", stat = "min") +
  ggttitle("Minimum: Observed vs. Predicted")

p4 <- pp_check(fit_rt, type = "stat", stat = "max") +
  ggttitle("Maximum: Observed vs. Predicted")

# Display all plots
library(patchwork)
(p1 | p2) / (p3 | p4)

```



## Extract and Analyze Posterior Predictions

```

# Draw from posterior predictive distribution
post_pred <- posterior_predict(fit_rt, ndraws = 1000)
dim(post_pred) # 1000 draws x n observations

```

[1] 1000 3000

```
cat("\nDimensions of posterior predictions:\n")
```

Dimensions of posterior predictions:

```
cat("Draws:", nrow(post_pred), "\n")
```

Draws: 1000

```
cat("Observations:", ncol(post_pred), "\n")
```


Observations: 3000

### Compare Observed vs. Predicted

```
# Compare observed vs. predicted means
obs_mean <- mean(rt_data$log_rt)
pred_mean <- apply(post_pred, 1, mean)

hist(pred_mean,
      main = "Posterior predictive distribution of mean log-RT",
      xlab = "Mean log-RT",
      col = "lightblue",
      breaks = 30)
abline(v = obs_mean, col = "red", lwd = 3, lty = 2)
legend("topright",
       legend = c("Predicted", "Observed"),
       col = c("lightblue", "red"),
       lwd = c(1, 3),
       lty = c(1, 2))
```

### Posterior predictive distribution of mean log-RT



```
cat("nObserved mean log-RT:", round(obs_mean, 3), "n")
```

Observed mean log-RT: 6.071

```
cat("Predicted mean log-RT (median):", round(median(pred_mean), 3), "n")
```

Predicted mean log-RT (median): 6.071

```
cat("95% CI for predicted mean:",
```

```
round(quantile(pred_mean, c(0.025, 0.975)), 3), "n")
```

95% CI for predicted mean: 6.056 6.087

## Check Posterior Predictive Intervals

```
# Check 95% posterior predictive interval
post_pred_interval <- apply(post_pred, 2, quantile, c(0.025, 0.975))

# Roughly 95% of observed values should fall within their interval
coverage <- mean(rt_data$log_rt > post_pred_interval[1,] &
                  rt_data$log_rt < post_pred_interval[2,])

cat("\nPosterior Predictive Interval Coverage:\n")
```

Posterior Predictive Interval Coverage:

```
cat("Proportion of observations within 95% interval:", round(coverage, 3), "\n")
```

Proportion of observations within 95% interval: 0.951

```
cat("Expected: ~0.95\n")
```

Expected: ~0.95

```
if (coverage < 0.90) {
  cat("\n Warning: Coverage is lower than expected.\n")
  cat("Model may be overconfident or missing important structure.\n")
} else if (coverage > 0.98) {
  cat("\n Warning: Coverage is higher than expected.\n")
  cat("Model may be too uncertain or overfitted.\n")
} else {
  cat("\n Coverage looks good!\n")
}
```

Coverage looks good!

## Check by Condition

```

# Group predictions by condition
rt_data_cond <- rt_data %>%
  group_by(condition) %>%
  summarise(
    obs_mean = mean(log_rt),
    obs_sd = sd(log_rt)
  )

# Get posterior predictions by condition
post_pred_A <- posterior_predict(fit_rt,
                                   newdata = filter(rt_data, condition == "A"),
                                   ndraws = 1000)
post_pred_B <- posterior_predict(fit_rt,
                                   newdata = filter(rt_data, condition == "B"),
                                   ndraws = 1000)


pred_mean_A <- apply(post_pred_A, 1, mean)
pred_mean_B <- apply(post_pred_B, 1, mean)

# Plot comparison
par(mfrow = c(1, 2))

hist(pred_mean_A,
      main = "Condition A: Mean log-RT",
      xlab = "Mean log-RT",
      col = "lightblue",
      breaks = 30,
      xlim = range(c(pred_mean_A, pred_mean_B)))
abline(v = rt_data_cond$obs_mean[1], col = "red", lwd = 3, lty = 2)

hist(pred_mean_B,
      main = "Condition B: Mean log-RT",
      xlab = "Mean log-RT",
      col = "lightgreen",
      breaks = 30,
      xlim = range(c(pred_mean_A, pred_mean_B)))
abline(v = rt_data_cond$obs_mean[2], col = "red", lwd = 3, lty = 2)

```



```
cat("\nBy Condition:\n")
```

By Condition:

```
cat("Condition A - Observed:", round(rt_data_cond$obs_mean[1], 3), "\n")
```

Condition A - Observed: 5.99

```
cat("Condition A - Predicted:", round(median(pred_mean_A), 3), "\n")
```

Condition A - Predicted: 5.99

```
cat("Condition B - Observed:", round(rt_data_cond$obs_mean[2], 3), "\n")
```

Condition B - Observed: 6.152

```
cat("Condition B - Predicted:", round(median(pred_mean_B), 3), "\n")
```

Condition B - Predicted: 6.151

## Summary

### Key Diagnostics Checked

1. **Visual inspection** - Observed data overlaps with posterior predictions
2. **Mean** - Central tendency captured correctly
3. **SD** - Spread of data captured correctly
4. **Extreme values** - Min/max are reasonable
5. **Predictive intervals** - Coverage is appropriate
6. **By condition** - Model captures group differences

### Common Problems and Solutions

Problem	Diagnosis	Solution
Model predictions too narrow	SD of posterior predictions < SD of data	Relax priors, check formula
Model predictions too wide	SD of posterior predictions » SD of data	Tighten priors, add more structure
Misses condition effects	Mean differs dramatically by condition	Add condition $\times$ random effect interaction
Extreme value mismatch	Min/max far from observed	Check for outliers, consider robust models

### Next Steps

If posterior predictive checks reveal problems:

1. **Adjust model formula** - Add missing predictors or interactions
2. **Revise priors** - May be too restrictive or too vague
3. **Consider alternative families** - E.g., Student's t for robust modeling
4. **Check for outliers** - May need to handle separately

```
sessionInfo()
```

```

R version 4.4.1 (2024-06-14)
Platform: x86_64-pc-linux-gnu
Running under: Ubuntu 22.04.5 LTS

Matrix products: default
BLAS: /usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3
LAPACK: /usr/lib/x86_64-linux-gnu/openblas-pthread/libopenblas-p0.3.20.so; LAPACK version 3.8.0

locale:
[1] LC_CTYPE=en_US.UTF-8          LC_NUMERIC=C
[3] LC_TIME=en_US.UTF-8          LC_COLLATE=en_US.UTF-8
[5] LC_MONETARY=en_US.UTF-8      LC_MESSAGES=en_US.UTF-8
[7] LC_PAPER=en_US.UTF-8         LC_NAME=C
[9] LC_ADDRESS=C                 LC_TELEPHONE=C
[11] LC_MEASUREMENT=en_US.UTF-8   LC_IDENTIFICATION=C

time zone: Etc/UTC
tzcode source: system (glibc)

attached base packages:
[1] stats      graphics    grDevices utils      datasets   methods    base

other attached packages:
[1] patchwork_1.3.2  bayesplot_1.14.0 lubridate_1.9.3 forcats_1.0.0
[5] stringr_1.5.1   dplyr_1.1.4    purrr_1.0.2    readr_2.1.5
[9] tidyverse_2.0.0  tibble_3.2.1   ggplot2_4.0.0  tidyverse_2.0.0
[13] brms_2.23.0    Rcpp_1.0.13

loaded via a namespace (and not attached):
[1] gtable_0.3.6        tensorA_0.36.2.1    xfun_0.54
[4] QuickJSR_1.8.1     inline_0.3.21      lattice_0.22-6
[7] tzdb_0.4.0          vctrs_0.6.5       tools_4.4.1
[10] generics_0.1.3     stats4_4.4.1      parallel_4.4.1
[13] fansi_1.0.6         pkgconfig_2.0.3    Matrix_1.7-0
[16] checkmate_2.3.3    RColorBrewer_1.1-3  S7_0.2.0
[19] distributional_0.5.0 RcppParallel_5.1.11-1 lifecycle_1.0.4
[22] compiler_4.4.1     farver_2.1.2      Brobdingnag_1.2-9
[25] tinytex_0.53        codetools_0.2-20   htmltools_0.5.8.1
[28] yaml_2.3.10        pillar_1.9.0      StanHeaders_2.32.10
[31] bridgesampling_1.1-2 abind_1.4-8      nlme_3.1-164
[34] posterior_1.6.1.9000 rstan_2.32.7    tidyselect_1.2.1
[37] digest_0.6.37      mvtnorm_1.3-3    stringi_1.8.4
[40] reshape2_1.4.4     labeling_0.4.3    fastmap_1.2.0

```

```
[43] grid_4.4.1           cli_3.6.5          magrittr_2.0.3
[46] loo_2.8.0            pkgbuild_1.4.8    utf8_1.2.4
[49] withr_3.0.2          scales_1.4.0        backports_1.5.0
[52] estimability_1.5.1   timechange_0.3.0   rmarkdown_2.30
[55] matrixStats_1.5.0    emmeans_2.0.0       gridExtra_2.3
[58] hms_1.1.3             coda_0.19-4.1      evaluate_1.0.1
[61] knitr_1.50            rstantools_2.5.0   rlang_1.1.6
[64] xtable_1.8-4          glue_1.8.0         jsonlite_1.8.9
[67] plyr_1.8.9            R6_2.5.1
```