

Bayes Factors and Hypothesis Testing

Bayesian Mixed Effects Models with brms for Linguists

Job Schepens

2026-01-06

Table of contents

1	The Question: Which Hypothesis is Better Supported?	2
1.1	Research Scenario	2
1.2	Bayes Factor vs. ROPE: Different Questions	3
1.3	Preview: Our Tools	3
2	Where We Are in the Analysis Workflow	3
2.1	The Bayesian Workflow So Far	3
2.2	LOO vs. Bayes Factors: What's the Difference?	4
3	The Savage-Dickey Density Ratio Method	4
3.1	What is the Savage-Dickey Method?	4
3.2	Visual Understanding	4
3.3	Why This Works (Technical)	5
4	Using hypothesis() for Bayes Factors	5
4.1	The hypothesis() Function	5
4.2	Example 1: Test for Effect Existence	5
4.2.1	Test: Does Complexity Increase Reading Time?	5
4.2.2	Directional Test	6
4.3	Example 2: Compare Two Groups	6
4.4	Example 3: Complex Contrasts	7
4.5	Critical Detail: sample_prior = "yes"	7
5	Interpreting Bayes Factors	8
5.1	Bayes Factor Scales	8
5.1.1	Jeffreys (1961) Scale	8
5.1.2	Lee & Wagenmakers (2013) - More Conservative	8
5.1.3	Visual Interpretation	8
5.2	What Bayes Factors Do NOT Tell You	8
5.3	Prior Sensitivity for Bayes Factors	9
6	Full Model Comparison with bayes_factor()	10
6.1	When hypothesis() is Not Enough	10
6.2	Bridge Sampling Method	10

6.3	Example: Compare Model Structures	10
6.4	Critical Detail: <code>save_pars = save_pars(all = TRUE)</code>	10
6.5	Comparing Multiple Models	11
6.6	Bridge Sampling Tips	11
7	Combining Approaches: A Complete Workflow	11
7.1	Recommended Analysis Strategy	11
7.2	Decision Matrix	11
7.3	Example: Complete Analysis	11
8	Reporting Bayes Factors in Papers	13
8.1	Methods Section Template	13
8.2	Results Section Template	14
8.3	Common Pitfalls to Avoid	14
9	Advanced Topics	15
9.1	Informed Priors from Previous Studies	15
9.2	Sequential Testing	15
9.3	Comparing More Than Two Hypotheses	15
10	Summary and Key Takeaways	15
10.1	What We Learned	15
10.2	When to Use What	15
10.3	Common Questions	16
10.4	Next Steps	16
11	Literature and Resources	16
11.1	Key Papers	16
11.1.1	Bayes Factors - Theory and Methods	16
11.1.2	Bayesian Hypothesis Testing - Applied	17
11.1.3	Integration with ROPE	17
11.1.4	Applications in Linguistics	17
11.2	Software Documentation	18
11.2.1	<code>brms</code>	18
11.2.2	<code>bridgesampling</code>	18
11.3	Online Resources	18
11.4	Module 06 (Previous)	18
11.5	Related Workshop Modules	19

1 The Question: Which Hypothesis is Better Supported?

1.1 Research Scenario

You’ve estimated your model parameters and found that an effect exists (Module 06). Now you want to answer a different question:

“How much evidence do the data provide for one hypothesis over another?”

Examples: - H : No effect ($= 0$) vs. H : Some effect ($\neq 0$) - H : Positive effect (> 0) vs. H : Negative effect (< 0) - Model A: Simple main effects vs. Model B: Include interaction

This is where **Bayes Factors** come in.

1.2 Bayes Factor vs. ROPE: Different Questions

Approach	Question	Output	Focus
ROPE (Module 06)	“Is the effect meaningful?”	Accept/Reject/Undecided	Practical significance
emmeans (Module 06)	“What are all pairwise differences?”	All comparisons, EMMs	Factorial designs
marginaleffects (Module 06)	“What’s the effect at X?”	Predictions, contrasts	Flexible estimation
Bayes Factor (Module 07)	“Which hypothesis is better?”	Evidence ratio (e.g., 10:1)	Hypothesis comparison

Key insight: These are complementary, not competing!

- Use ROPE when you care about **practical significance**
- Use emmeans/marginaleffects for **effect estimation** in complex designs
- Use Bayes Factors when comparing **competing hypotheses/theories**

1.3 Preview: Our Tools

Today we’ll learn:

1. **hypothesis() function:** Compute Bayes Factors for parameter constraints
 - Uses Savage-Dickey density ratio method
 - Fast and built into brms
2. **bayes_factor() function:** Compare full models
 - Uses bridge sampling (more general but slower)
 - For complex model comparisons
3. **Interpretation guidelines:** What does $BF = 3$ vs. $BF = 30$ mean?

2 Where We Are in the Analysis Workflow

2.1 The Bayesian Workflow So Far

Module 01-02: Build model + set priors

↓

Module 03: Check model fit (posterior predictive)

↓

Module 04: Test prior sensitivity

↓

Module 05: Compare models (LOO-CV for prediction)

↓

Module 06: Practical significance (ROPE, emmeans, marginaleffects)

↓

Module 07 (TODAY): Hypothesis comparison (Bayes Factors, hypothesis())

2.2 LOO vs. Bayes Factors: What's the Difference?

Both compare models, but with different goals:

LOO (Module 05): - Goal: Predictive accuracy - Question: "Which model predicts new data better?"
- Method: Cross-validation - Use when: You care about out-of-sample performance - Example: Choose between polynomial degrees for smooth fit

Bayes Factors (Module 07): - Goal: Relative evidence - Question: "Which model is better supported by data?" - Method: Ratio of marginal likelihoods - Use when: You want to quantify evidence for theory - Example: Does priming effect exist? (H vs. H)

Rule of thumb: - Use LOO for model selection when **prediction** matters - Use BF for model comparison when **explanation/theory** matters

3 The Savage-Dickey Density Ratio Method

3.1 What is the Savage-Dickey Method?

The **Savage-Dickey method** provides an elegant way to compute Bayes Factors when: - Models are **nested** (one is a special case of the other) - You're testing a point hypothesis (e.g., $\theta = 0$)

The formula:

$$BF_{01} = \frac{p(\theta = \theta_0 | \text{Data}, H_1)}{p(\theta = \theta_0 | H_1)} = \frac{\text{posterior density at null}}{\text{prior density at null}}$$

Intuition: - If data make null value **more plausible** $\rightarrow BF > 1 \rightarrow$ evidence for H - If data make null value **less plausible** $\rightarrow BF < 1 \rightarrow$ evidence for H

3.2 Visual Understanding

Savage-Dickey Density Ratio

Data shifted belief away from null $\rightarrow$ Evidence for H.




Figure 1: Savage-Dickey Density Ratio Illustration

Interpretation of this example: - Prior density at $\theta = 0$: 0.402 - Posterior density at $\theta = 0$: 0.356 - The posterior density **decreased** at the null - $BF = 1.13 \rightarrow$ Data are ~1.1 times more likely under H_1 - This is **positive evidence** for an effect

3.3 Why This Works (Technical)

For nested models where $H_0: \theta = 0$ is a special case of $H_1: \theta \neq 0$:

$$BF_{01} = \frac{p(\text{Data}|H_0)}{p(\text{Data}|H_1)}$$

The Savage-Dickey method shows that this equals:

$$BF_{01} = \frac{p(\theta = 0|\text{Data}, H_1)}{p(\theta = 0|H_1)}$$

This is exact, not an approximation! (under specific conditions)

Conditions required: 1. H_0 is a special case of H_1 (nested models) 2. Prior on nuisance parameters is same in both models 3. You can accurately estimate densities at the null value

4 Using hypothesis() for Bayes Factors

4.1 The hypothesis() Function

The hypothesis() function in brms implements the Savage-Dickey method automatically.

Basic syntax:

```
hypothesis(model, hypothesis = "parameter = value")
hypothesis(model, hypothesis = "parameter > value")
hypothesis(model, hypothesis = "parameter1 - parameter2 = 0")
```

4.2 Example 1: Test for Effect Existence

4.2.1 Test: Does Complexity Increase Reading Time?

Hypothesis Tests for class b:

	Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio
1	(complexitySimple) = 0	-0.12	0.01	-0.13	-0.1	0
	Post.Prob	Star				
1		0				*

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

Understanding the output:

```

Hypothesis: complexitySimple = 0
Estimate: -0.08
CI.Lower: -0.11
CI.Upper: -0.05
Evid.Ratio: 0.02
Post.Prob: 0.02
Star: *

```

Interpretation: - **Estimate:** Complex sentences are 0.08 log-units slower - **95% CI:** [-0.11, -0.05] (doesn't include 0) - **Evid.Ratio:** 0.02 means $BF = 0.02$, so $BF = 1/0.02 = 50$ - **Conclusion:** Data are **50 times more likely** under H (effect exists) than H (no effect) - This is **very strong evidence** for a complexity effect

4.2.2 Directional Test

Hypothesis Tests for class b:

Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio	Post.Prob	Star
1 (complexitySimple) < 0	-0.12	0.01	-0.13	-0.1	Inf		
1	1						*

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

Why directional tests? - Theory predicts direction → stronger evidence possible - One-sided test has more power than two-sided - Evidence Ratio will be higher if data support predicted direction

4.3 Example 2: Compare Two Groups

Hypothesis Tests for class b:

Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio	Post.Prob	Star
1 (groupNative) = 0	0.25	0.23	-0.19	0.69	1.16	0.54	
1							

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

Proportion of samples inside the ROPE [-0.20, 0.20]:

Parameter	Inside ROPE
Intercept	0.00 %
groupNative	40.97 %

Integrated interpretation: - **Bayes Factor:** Quantifies evidence for difference - **ROPE:** Checks if

difference is meaningful - **Both together**: Complete picture

4.4 Example 3: Complex Contrasts

Hypothesis Tests for class b:

	Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio
1	(groupNative:cond... = 0	-0.1	0.05	-0.2	0	0.57
	Post.Prob	Star				
1		0.36	*			

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

Hypothesis Tests for class b:

	Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio
1	(groupNative:cond... = 0	-0.38	0.11	-0.6	-0.16	0.02
	Post.Prob	Star				
1		0.02	*			

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

Hypothesis Tests for class b:

	Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio
1	(Intercept)-(5.5) > 0	0.50	0.03	0.45	0.56	Inf
2	(conditionB) > 0	0.14	0.04	0.08	0.20	Inf
3	(groupNative:cond... > 0	-0.10	0.05	-0.18	-0.02	0.02
	Post.Prob	Star				
1		1.00	*			
2		1.00	*			
3		0.02				

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

Advanced hypothesis testing: - Test interaction terms directly - Test specific numerical relationships (e.g., "effect is twice as large") - Test multiple hypotheses simultaneously

4.5 Critical Detail: `sample_prior = "yes"`

You **MUST** include `sample_prior = "yes"` in `brm()` to use `hypothesis()`!

Why? - Savage-Dickey needs both prior and posterior samples - By default, brms only samples the posterior

- `sample_prior = "yes"` also samples from the prior distribution

```
# WRONG - will fail
model_no_prior <- brm(..., sample_prior = "no")
hypothesis(model_no_prior, "b = 0") # ERROR!

# CORRECT
model_with_prior <- brm(..., sample_prior = "yes")
hypothesis(model_with_prior, "b = 0") # Works!
```

Performance note: This adds minimal computational cost (~5% overhead).

5 Interpreting Bayes Factors

5.1 Bayes Factor Scales

Different researchers use different interpretation scales. Here are three common ones:

5.1.1 Jeffreys (1961) Scale

BF	Evidence for H
> 100	Decisive
30-100	Very strong
10-30	Strong
3-10	Substantial
1-3	Anecdotal
1	No evidence
1/3-1	Anecdotal for H
1/10-1/3	Substantial for H
< 1/10	Strong for H

5.1.2 Lee & Wagenmakers (2013) - More Conservative

BF	Evidence for H
> 100	Extreme
30-100	Very strong
10-30	Strong
3-10	Moderate
1-3	Anecdotal

5.1.3 Visual Interpretation

5.2 What Bayes Factors Do NOT Tell You

Common misconceptions:

1. **BF = probability of hypothesis being true**
 - $BF = 10$ does NOT mean “90% probability H is true”

Bayes Factor Interpretation Scale

BF.: Evidence for H. relative to H.




Figure 2: Interpreting Bayes Factors

- BF quantifies relative evidence, not absolute probability
2. **BF depends on prior specification**
 - Vague priors → built-in Ockham's razor → penalize complex models
 - This is a feature, not a bug!
 - Always report your priors
 3. **BF is continuous, not a decision rule**
 - Don't treat $BF > 3$ as "threshold for significance"
 - Report the actual BF value
 - Let readers judge based on context
 4. **BF doesn't tell you effect size**
 - $BF = 100$ could mean tiny effect with large N
 - Always report effect estimates + uncertainty
 - Combine with ROPE analysis (Module 06)

5.3 Prior Sensitivity for Bayes Factors

Bayes Factors are **more sensitive to priors** than posterior estimates.

Demonstration:

Vague prior BF : 1.48571e+31

Informative prior BF : 1.881926e+28

Why they differ: - Vague prior spreads probability mass widely - Data must "overcome" this large prior space - Informative prior concentrates mass where effect likely is - Same posterior, different BF!

Best practice: - Use realistic informative priors based on domain knowledge - Report prior specification in paper - Consider prior sensitivity analysis

6 Full Model Comparison with `bayes_factor()`

6.1 When `hypothesis()` is Not Enough

The `hypothesis()` function (Savage-Dickey) works for: - Nested models - Point hypotheses on parameters

But what if you want to compare: - **Non-nested models** (e.g., Model A: $y \sim x_1$ vs. Model B: $y \sim x_2$) - **Complex model structures** (different random effects) - **Different families** (Gaussian vs. Student-t)

For these cases, use `bayes_factor()` with **bridge sampling**.

6.2 Bridge Sampling Method

Bridge sampling estimates the marginal likelihood:

$$p(\text{Data}|H) = \int p(\text{Data}|\theta, H) \cdot p(\theta|H) d\theta$$

Then computes:

$$BF_{12} = \frac{p(\text{Data}|H_1)}{p(\text{Data}|H_2)}$$

Advantages: - Works for any model comparison - Doesn't require nesting

Disadvantages: - Computationally intensive - Requires additional package: `bridgesampling`

6.3 Example: Compare Model Structures

```
Loaded cached marginal likelihood for RI model
```

```
Loaded cached marginal likelihood for RS model
```

```
Estimated Bayes factor in favor of x1 over x2: 0.50648
```

Interpretation:

```
Estimated Bayes factor in favor of ml_rs over ml_ri: 3.45
```

- $BF = 3.45 \rightarrow$ Moderate evidence for random slopes model
- Random slopes provide better explanation of data
- But: Also check LOO (Module 05) for predictive accuracy!

6.4 Critical Detail: `save_pars = save_pars(all = TRUE)`

For bridge sampling, you **MUST** save all parameters!

```
# WRONG - bridge sampling will fail
model <- brm(..., save_pars = save_pars(all = FALSE))

# CORRECT
model <- brm(..., save_pars = save_pars(all = TRUE))
```

Trade-off: - Saves all parameters $\rightarrow$ larger model objects - But: Necessary for bridge sampling to work - Only include this when you plan to use `bayes_factor()`

6.5 Comparing Multiple Models

Comparison	BF	Evidence
1 RI vs Null	1.315199e+38	For complexity effect
2 RS vs Null	6.661226e+37	For complexity + random slopes
3 RS vs RI	5.064805e-01	For adding random slopes

6.6 Bridge Sampling Tips

1. Use enough iterations: - Minimum: 4000 iterations (2000 warmup) - Better: 10000+ iterations for stable estimates

2. Check convergence:

```
# Bridge sampling has its own convergence diagnostic
print(ml_ri) # Look for "relative mean-squared error" (should be small)
```

3. Combine with LOO: - BF → Which model explains data better? - LOO → Which model predicts better? - Use both for complete picture

4. Computational cost: - Bridge sampling is slow for complex models - Consider using `hypothesis()` when possible (much faster)

7 Combining Approaches: A Complete Workflow

7.1 Recommended Analysis Strategy

7.2 Decision Matrix

Bayes Factor	ROPE Result (Module 06)	Interpretation
BF > 10	Outside ROPE	Strong evidence for meaningful effect
BF > 10	Inside ROPE	Effect exists but too small to matter
BF 1-10	Outside ROPE	Meaningful effect but moderate evidence
BF > 10	Inside ROPE	Strong evidence effect is negligible
BF 1	Overlaps ROPE	Undecided - collect more data

How to use this table: 1. Compute Bayes Factor (Module 07) → strength of evidence 2. Check ROPE (Module 06) → practical significance 3. Use `emmeans/marginalEffects` (Module 06) → effect estimation

7.3 Example: Complete Analysis

=== Main Effect Test ===

Hypothesis Tests for class b:

	Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio	Post.Prob
1	(groupNative) = 0	-0.01	0.01	-0.03	0	7.16	0.88
	Star						

1

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.

=== ROPE Analysis ===

Proportion of samples inside the ROPE [-0.05, 0.05]:

Parameter	Inside ROPE
Intercept	0.00 %
groupNative	100.00 %
clause_typeSubject	0.00 %
groupNative:clause_typeSubject	100.00 %

=== Interaction Test ===

Hypothesis Tests for class b:

	Hypothesis	Estimate	Est.Error	CI.Lower	CI.Upper	Evid.Ratio
1	(groupNative:clau... = 0	-0.01	0.01	-0.03	0.02	13.4
	Post.Prob Star					


1 0.93

'CI': 90%-CI for one-sided and 95%-CI for two-sided hypotheses.

'*': For one-sided hypotheses, the posterior probability exceeds 95%;

for two-sided hypotheses, the value tested against lies outside the 95%-CI.

Posterior probabilities of point hypotheses assume equal prior probabilities.



Integrated interpretation:

1. **Effect size:** Report posterior mean + 95% CI
2. **Practical significance:** Is 95% HDI outside ROPE?
3. **Evidence strength:** What's the Bayes Factor?
4. **Conclusion:** Combine all three pieces

8 Reporting Bayes Factors in Papers

8.1 Methods Section Template

Hypothesis Testing

We quantified evidence for hypotheses using Bayes Factors computed via the Savage-Dickey density ratio method (Wagenmakers et al., 2010), implemented in the `brms::hypothesis()` function (Bürkner, 2017).

Bayes Factors express the relative evidence for one hypothesis over another as a ratio of marginal likelihoods. We interpret $BF > 3$ as

positive evidence, $BF > 10$ as strong evidence, and $BF > 30$ as very strong evidence for the alternative hypothesis (Lee & Wagenmakers, 2013).

Priors for regression coefficients were specified as $\text{Normal}(0, 0.2)$ on the log-RT scale, reflecting realistic effect sizes for psycholinguistic experiments (Nicenboim et al., 2023). We conducted sensitivity analyses with alternative prior specifications (see Supplementary Materials).

For practical significance testing, we used the ROPE framework (Kruschke, 2018) with `bayestestR` (Makowski et al., 2019). Effect estimation for factorial designs was conducted using `emmeans` (Lenth, 2016) and `marginaleffects` (Arel-Bundock et al., 2024).

8.2 Results Section Template

Reading Time Analysis

Complex sentences elicited longer reading times than simple sentences ($\beta = 0.08$, 95% CI: [0.05, 0.11]), corresponding to a median increase of 8% on the millisecond scale.

There was very strong evidence for this effect ($BF = 45.3$), indicating the data were approximately 45 times more likely under the hypothesis of an effect than under the null hypothesis of no effect.

Using a ROPE of ± 0.03 log-units (corresponding to $\pm 3\%$ change in reading time, which we consider the threshold for practical significance; Kruschke, 2018), we found that the entire 95% credible interval excluded the ROPE, indicating the effect was both statistically credible and practically meaningful.

Pairwise comparisons using `emmeans` confirmed that all three conditions differed meaningfully from each other (all 95% CIs excluded the ROPE), with average differences ranging from 0.08 to 0.15 log-units.

8.3 Common Pitfalls to Avoid

Don't say: - “ $BF = 10$, therefore the effect is real with 90% probability” - “ $BF > 3$, so we reject the null hypothesis” - “ $BF = 2.5$, which is not significant”

Do say: - “ $BF = 10$, indicating strong evidence for H_1 relative to H_0 ” - “The data are 10 times more likely under the alternative hypothesis” - “ $BF = 2.5$ provides weak to moderate evidence”

Always include: 1. Prior specification (what priors you used) 2. Interpretation scale (which guidelines you follow) 3. Effect size + uncertainty (BF alone is not enough) 4. Context (what does this evidence mean for your research question?)

9 Advanced Topics

9.1 Informed Priors from Previous Studies

Informed prior BF : 1.057572e+23

Default prior BF : 3.315914e+28

When to use informed priors: - You have strong theoretical predictions - Previous literature provides effect size estimates - You want to show robustness to prior specification

Transparency: - Always report both analyses (informed and default) - Justify your informed prior with citations - Show prior sensitivity analysis

9.2 Sequential Testing

Unlike frequentist p-values, Bayes Factors do **not** suffer from multiple testing problems!

Advantages: - Stop data collection when evidence is conclusive - More efficient than fixed-N designs - Ethically appropriate (don't over-collect data)

9.3 Comparing More Than Two Hypotheses

10 Summary and Key Takeaways

10.1 What We Learned

1. **Bayes Factors quantify evidence** for hypotheses
 - Not probability of hypothesis being true
 - Ratio of evidence between two hypotheses
2. **Savage-Dickey method** (via `hypothesis()`)
 - Fast and easy for nested models
 - Tests point hypotheses on parameters
 - Requires `sample_prior = "yes"`
3. **Bridge sampling** (via `bayes_factor()`)
 - Works for any model comparison
 - Computationally expensive
 - Requires `save_pars = save_pars(all = TRUE)`
4. **Interpretation guidelines**
 - $BF > 10$ = strong evidence
 - $BF \text{ } 3\text{-}10$ = moderate evidence
 - $BF < 3$ = weak/anecdotal evidence
5. **Combine with Module 06 tools**
 - BF (Module 07) → strength of evidence for hypotheses
 - ROPE (Module 06) → practical significance
 - `emmeans/marginalEffects` (Module 06) → effect estimation
 - Together → complete picture

10.2 When to Use What

Use `hypothesis()` when: - Testing specific parameter values ($= 0$, > 0) - Models are nested - You want fast computation - Testing directional predictions

Use `bayes_factor()` when: - Comparing non-nested models - Different model structures - Need comprehensive model comparison

Use ROPE (Module 06) when: - Focus on practical significance - Want accept/reject/undecided decision rule - Testing if effect is negligible

Use `emmeans` (Module 06) when: - Factorial designs - Need all pairwise comparisons - Familiar with traditional EMM workflow

Use `marginalEffects` (Module 06) when: - Flexible predictions at specific values - Custom contrasts and comparisons - Working with continuous predictors

Use LOO (Module 05) when: - Focus on prediction accuracy - Model selection for forecasting - Cross-validation needed

10.3 Common Questions

Q: Should I always report Bayes Factors?

A: No. Report BF when: - Your research question is about comparing specific hypotheses - You want to quantify strength of evidence - You're testing theoretically-motivated predictions

Don't report BF if: - Your focus is purely exploratory - You're mainly interested in effect size estimation - BF would be redundant with ROPE analysis

Q: What prior should I use for BF?

A: Use **weakly informative priors** based on domain knowledge: - Review previous literature for effect size estimates - Consider measurement scale (log-RT vs. RT) - Use prior predictive checks (Module 02) - Report sensitivity analysis

Q: Can I use BF for model selection?

A: Yes, but: - Combine with LOO for prediction assessment - BF favors explanation, LOO favors prediction - Use both when possible - Report effect sizes regardless

10.4 Next Steps

Practice exercises: 1. Compute BF for directional hypothesis in your data 2. Compare two models with different random effects structures 3. Conduct prior sensitivity analysis for BF 4. Create integrated report with ROPE + BF + effect sizes

Further reading: - Wagenmakers et al. (2010) - Savage-Dickey method - Kass & Raftery (1995) - Bayes Factors overview - van Ravenzwaaij & Wagenmakers (2022) - Advantages of Bayes - Schad et al. (2024) - Workflow for linguistic data

11 Literature and Resources

11.1 Key Papers

11.1.1 Bayes Factors - Theory and Methods

- Wagenmakers, E.-J., Lodewyckx, T., Kuriyal, H., & Grasman, R. (2010). Bayesian hypothesis testing for psychologists: A tutorial on the Savage–Dickey method. *Cognitive Psychology*,

60(3), 158-189.

- **Essential reading** for understanding Savage-Dickey method
- Complete worked examples with code
- Connection to `brms::hypothesis()`
- **Packages:** Foundational theory for `brms::hypothesis()`
- **Kass, R. E., & Raftery, A. E. (1995).** Bayes factors. *Journal of the American Statistical Association*, 90(430), 773-795.
 - Classic reference on Bayes Factors
 - Mathematical foundations
 - Interpretation guidelines
 - **Packages:** Conceptual (applicable to all Bayesian software)
- **Gronau, Q. F., Sarafoglou, A., Matzke, D., Ly, A., Boehm, U., Marsman, M., ... & Steingroever, H. (2017).** A tutorial on bridge sampling. *Journal of Mathematical Psychology*, 81, 80-97.
 - Bridge sampling method explained
 - Computational details
 - **Packages:** `bridgesampling`, used with `brms`

11.1.2 Bayesian Hypothesis Testing - Applied

- **van Ravenzwaaij, D., & Wagenmakers, E.-J. (2022).** Advantages masquerading as ‘issues’ in Bayesian hypothesis testing: A commentary on Tendeiro and Kiers (2019). *Psychological Methods*, 27(3), 451-465.
 - Addresses common criticisms of Bayesian methods
 - Philosophy of hypothesis testing
 - **Packages:** Conceptual (defends Bayesian approach generally)
- **Lee, M. D., & Wagenmakers, E.-J. (2013).** *Bayesian cognitive modeling: A practical course*. Cambridge University Press.
 - Practical guide to Bayesian hypothesis testing
 - Many worked examples
 - Interpretation scales for BF
 - **Packages:** `JAGS`, `WinBUGS` (concepts transfer to `brms`)

11.1.3 Integration with ROPE

- **Kruschke, J. K. (2018).** Rejecting or accepting parameter values in Bayesian estimation. *Advances in Methods and Practices in Psychological Science*, 1(2), 270-280.
 - HDI + ROPE decision rule
 - Comparison with Bayes Factors
 - When to use each approach
 - **Packages:** `bayestestR::rope()`, works with `brms`

11.1.4 Applications in Linguistics

- **Nicenboim, B., Schad, D. J., & Vasisht, S. (2023).** *An introduction to Bayesian data analysis for cognitive science*.
 - Chapter on hypothesis testing with `brms`
 - Linguistic examples
 - Prior specification guidance

- **Packages:** `brms`, `hypothesis()`, complete workflows
- **Schad, D. J., Nicenboim, B., Bürkner, P.-C., Betancourt, M., & Vasishth, S. (2024).** Workflow techniques for the robust use of Bayes factors. *Psychological Methods*.
 - Best practices for BF in practice
 - Sensitivity analysis
 - Reporting guidelines
 - **Packages:** `brms`, `bridgesampling`, `hypothesis()`

11.2 Software Documentation

11.2.1 `brms`

- **`hypothesis()` documentation:** <https://paulbuerkner.com/brms/reference/hypothesis.brmsfit.html>
 - **`brms::hypothesis()`** complete API
 - Syntax for complex hypotheses
 - Examples of directional tests
- **Bürkner, P.-C. (2017).** `brms`: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software*, 80(1), 1-28.
 - Primary citation for `brms`
 - Technical details
 - **`brms`** package paper

11.2.2 `bridgesampling`

- **`bridgesampling` documentation:** <https://cran.r-project.org/package=bridgesampling>
 - **`bridgesampling`** package reference
 - Integration with **`brms`**
 - **`bayes_factor()`** function

11.3 Online Resources

- **Paul Bürkner's Blog:** <https://paulbuerkner.com/software/brms-blogposts.html>
 - Community posts on **`brms`**
 - `hypothesis()` examples
- **Stan Discourse Forum:** <https://discourse.mc-stan.org/>
 - Active community
 - Bayes Factor discussions
 - **`brms`** technical support

11.4 Module 06 (Previous)

- **Practical Significance and Effect Estimation:** See Module 06 for:
 - ROPE framework (`bayestestR` package)
 - Effect estimation with `emmeans` (factorial designs)
 - Flexible predictions with `marginalEffects`
 - Combining ROPE with `emmeans`/`marginalEffects`
 - **`bayestestR`, `emmeans`, `marginalEffects`**

11.5 Related Workshop Modules

- Module 05: Model comparison with LOO (predictive accuracy)
- Module 06: Practical significance (ROPE, emmeans, marginaeffects)
- Module 07 (this): Hypothesis comparison (Bayes Factors, hypothesis())