

Prior Predictive Checks: Reaction Time Example

Bayesian Mixed Effects Models with brms for Linguists

Workshop Materials

2025-11-11

Table of contents

| | | |
|----------|---|----------|
| 1 | Prior Predictive Checks for Reaction Time Data | 1 |
| 1.1 | Setup | 2 |
| 1.2 | Fitting Prior Only Model | 2 |
| 1.3 | Prior Distributions | 8 |
| 1.3.1 | Intercept Prior | 8 |
| 1.3.2 | Effect Size Prior | 10 |
| 1.3.3 | Residual Noise Prior | 11 |
| 1.4 | Random Effects Distributions | 11 |
| 1.4.1 | Subject Random Intercepts | 11 |
| 1.4.2 | Subject Random Slopes | 13 |
| 1.4.3 | Residual Noise Distribution | 15 |
| 1.4.4 | Random Effects Comparison | 16 |
| 1.5 | Interpretation | 17 |
| 1.5.1 | Good Signs (Prior is Reasonable) | 17 |
| 1.5.2 | Problems to Watch For | 17 |
| 1.6 | Summary | 17 |

1 Prior Predictive Checks for Reaction Time Data

This document demonstrates how to validate priors for a Bayesian RT model before fitting to data.

1.1 Setup

```
library(brms)
library(tidyverse)
library(bayesplot)
library(posterior) # For as_draws_df() - following Kurz's approach

# Create example RT data
set.seed(42)
n_subj <- 20
n_trials <- 50
n_items <- 30

rt_data <- expand.grid(
  trial = 1:n_trials,
  subject = 1:n_subj,
  item = 1:n_items
) %>%
  filter(row_number() <= n_subj * n_trials * 3) %>%
  mutate(
    condition = rep(c("A", "B"), length.out = n()),
    log_rt = rnorm(n(), mean = 6, sd = 0.3) +
      (condition == "B") * 0.15 +
      rnorm(n(), mean = 0, sd = 0.1),
    rt = exp(log_rt)
  )

# Define priors
rt_priors <- c(
  prior(normal(6, 1.5), class = Intercept),
  prior(normal(0, 0.5), class = b),
  prior(exponential(1), class = sigma),
  prior(exponential(1), class = sd),
  prior(lkj(2), class = cor)
)
```

1.2 Fitting Prior Only Model

```
# Create fits directory if it doesn't exist
if (!dir.exists("fits")) dir.create("fits")
```

```

# Fit model with priors only
# Using file argument to cache the fitted model and speed up re-runs
prior_pred <- brm(
  log_rt ~ condition + (1 + condition | subject) + (1 | item),
  data = rt_data,
  family = gaussian(),
  prior = rt_priors,
  sample_prior = "only",
  chains = 4, iter = 1000, # 4 chains, fewer iterations for prior checks
  cores = 4, # Use 4 cores for parallel sampling
  verbose = FALSE,
  refresh = 0,
  file = "fits/prior_pred_rt", # Cache model - only refits if data/formula/priors change
  file_refit = "on_change" # Only refit when necessary
)

```


```
## Prior Predictive Checks
```

```
### Visual Checks
```

```
::: {.cell}
```

```
```{.r .cell-code}
Density overlay
pp_check(prior_pred, type = "dens_overlay", ndraws = 100, prefix = "ppd") +
 labs(title = "Prior Predictions vs Observed Data (Density Overlay)")
```

## Prior Predictions vs Observed Data (Density Overlay)




```
Mean comparison
pp_check(prior_pred, type = "stat", stat = "mean") +
 labs(title = "Prior Predictions vs Observed Data (Mean)")
```

Using all posterior draws for ppc type 'stat' by default.

Note: in most cases the default test statistic 'mean' is too weak to detect anything of interest.

`stat\_bin()` using `bins = 30`. Pick better value `binwidth`.


## Prior Predictions vs Observed Data (Mean)



```
SD comparison
pp_check(prior_pred, type = "stat", stat = "sd") +
 labs(title = "Prior Predictions vs Observed Data (SD)")
```

Using all posterior draws for ppc type 'stat' by default.  
`stat\_bin()` using `bins = 30`. Pick better value `binwidth`.


## Prior Predictions vs Observed Data (SD)



```
Min comparison
pp_check(prior_pred, type = "stat", stat = "min") +
 labs(title = "Prior Predictions vs Observed Data (Min)")
```

Using all posterior draws for ppc type 'stat' by default.  
`stat\_bin()` using `bins = 30`. Pick better value `binwidth`.


## Prior Predictions vs Observed Data (Min)



```
Max comparison
pp_check(prior_pred, type = "stat", stat = "max") +
 labs(title = "Prior Predictions vs Observed Data (Max)")
```

Using all posterior draws for ppc type 'stat' by default.  
`stat\_bin()` using `bins = 30`. Pick better value `binwidth`.

## Prior Predictions vs Observed Data (Max)



...:

### 1.3 Prior Distributions

Following Kurz's approach, we extract prior samples using `as_draws_df()` from the `posterior` package.

```
Extract prior draws as a data frame (Kurz's method)
This works seamlessly with tidyverse and gives us a proper data frame
library(posterior)
prior_samples <- as_draws_df(prior_pred)
```

#### 1.3.1 Intercept Prior

```
Now we can access columns directly from the data frame
intercept_vals <- prior_samples$b_Intercept

cat("Intercept prior (log scale):\n")
```

Intercept prior (log scale):

```
print(quantile(intercept_vals, c(0.025, 0.5, 0.975), na.rm = TRUE))
```

```
2.5% 50% 97.5%
2.875128 6.020023 8.923694
```

```
cat("\nIntercept prior (RT scale in milliseconds):\n")
```

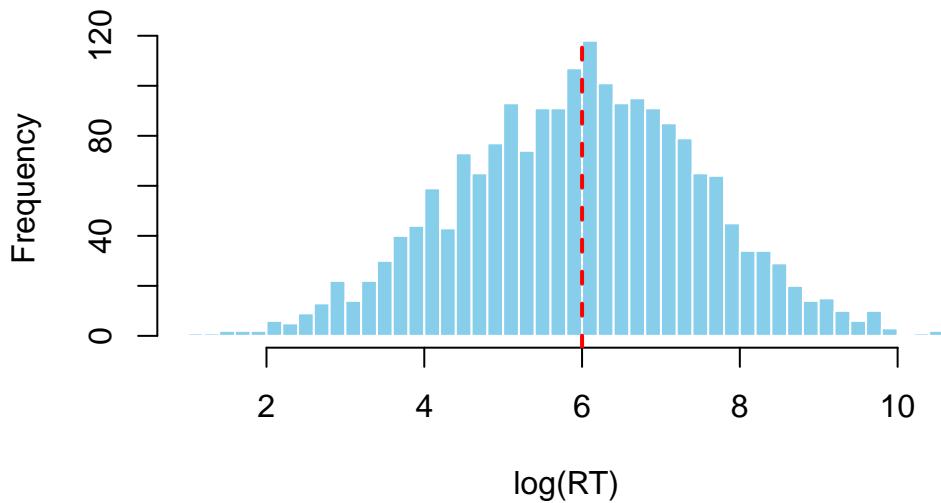
```
Intercept prior (RT scale in milliseconds):
```

```
print(exp(quantile(intercept_vals, c(0.025, 0.5, 0.975), na.rm = TRUE)))
```

```
2.5% 50% 97.5%
17.72769 411.58799 7507.77385
```

```
Visualization
hist(intercept_vals,
 main = "Prior for Intercept (log-RT scale)",
 xlab = "log(RT)",
 breaks = 50, col = "skyblue", border = "white")
abline(v = 6, col = "red", lwd = 2, lty = 2)
```

## Prior for Intercept (log-RT scale)



### 1.3.2 Effect Size Prior

```
effect_vals <- prior_samples$b_conditionB
cat("Condition effect prior (log scale):\n")
```

Condition effect prior (log scale):

```
effect_q <- quantile(effect_vals, c(0.025, 0.5, 0.975), na.rm = TRUE)
print(effect_q)
```

2.5% 50% 97.5%  
-0.99889479 -0.01414723 0.94791795

```
cat("\nCondition effect prior (RT scale in ms):\n")
```


Condition effect prior (RT scale in ms):

```
print(exp(effect_q))
```

2.5% 50% 97.5%  
0.3682863 0.9859524 2.5803317

```
Visualization
hist(effect_vals,
 main = "Prior for Condition Effect (log-RT scale)",
 xlab = "Condition B effect",
 breaks = 50, col = "lightcoral", border = "white")
abline(v = 0, col = "red", lwd = 2, lty = 2)
```

## Prior for Condition Effect (log-RT scale)



### 1.3.3 Residual Noise Prior

```
sigma_vals <- prior_samples$sigma
cat("Residual noise prior (sigma, log scale):\n")

Residual noise prior (sigma, log scale):

sigma_q <- quantile(sigma_vals, c(0.025, 0.5, 0.975), na.rm = TRUE)
print(sigma_q)
```

2.5%            50%            97.5%  
0.02573465 0.69876976 3.81711126

## 1.4 Random Effects Distributions

For prior predictive checks, we examine the **implied distribution** of subject-specific parameters by extracting the hyperprior SDs and simulating random effects.

### 1.4.1 Subject Random Intercepts

```
Extract hyperprior SDs from prior samples
sd_subject_intercept <- prior_samples$sd_subject_Intercept

cat("Subject random intercept SD prior:\n")
```

Subject random intercept SD prior:

```
print(quantile(sd_subject_intercept, c(0.025, 0.5, 0.975)))
```

```
2.5% 50% 97.5%
0.02519548 0.67840017 3.61983168
```

```
Simulate random effects from this prior
set.seed(123)
n_sims <- 1000
simulated_intercepts <- rnorm(n_sims, mean = 0, sd = median(sd_subject_intercept))

cat("\nImplied subject random intercepts (log scale):\n")
```

Implied subject random intercepts (log scale):

```
subject_int <- quantile(simulated_intercepts, c(0.025, 0.5, 0.975))
print(subject_int)
```

```
2.5% 50% 97.5%
-1.317150788 0.006247821 1.382502744
```

```
cat("\nImplied subject-specific RTs (milliseconds):\n")
```

```
Combine with typical intercept value from prior
intercept_median <- median(intercept_vals)
subject_rt <- exp(intercept_median + subject_int)
print(subject_rt)
```

```

2.5% 50% 97.5%
110.2634 414.1676 1640.1214

```

```

cat("Prior implies subject RTs range from",
 round(subject_rt[1]), "to", round(subject_rt[3]), "ms\n")

```

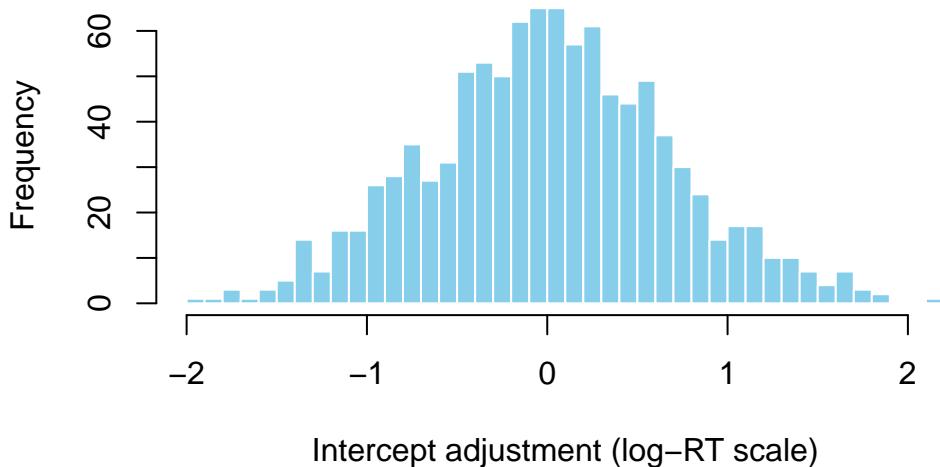
Prior implies subject RTs range from 110 to 1640 ms

```

Visualization
hist(simulated_intercepts,
 main = "Prior-implied Subject Random Intercepts",
 xlab = "Intercept adjustment (log-RT scale)",
 breaks = 30, col = "skyblue", border = "white")

```

## Prior-implied Subject Random Intercepts



### 1.4.2 Subject Random Slopes

```

Extract hyperprior SDs for slopes
sd_subject_slope <- prior_samples$sd_subject__conditionB

cat("Subject random slope SD prior:\n")

```

Subject random slope SD prior:

```
print(quantile(sd_subject_slope, c(0.025, 0.5, 0.975)))
```

```
2.5% 50% 97.5%
0.0249993 0.6998014 3.6529919
```

```
Simulate random slope effects
simulated_slopes <- rnorm(n_sims, mean = 0, sd = median(sd_subject_slope))

cat("\nImplied subject random slopes (log scale):\n")
```

Implied subject random slopes (log scale):

```
subject_slope <- quantile(simulated_slopes, c(0.025, 0.5, 0.975))
print(subject_slope)
```

```
2.5% 50% 97.5%
-1.39369961 0.03838577 1.33375954
```

```
cat("\nInterpretation: Condition effect varies by subject\n")
```

Interpretation: Condition effect varies by subject

```
cat("Small effect subjects (2.5%): ", round(exp(subject_slope[1]), 3), "x multiplier\n")
```

Small effect subjects (2.5%): 0.248 × multiplier


```
cat("Average effect subjects (50%): ", round(exp(subject_slope[2]), 3), "x multiplier\n")
```

Average effect subjects (50%): 1.039 × multiplier

```
cat("Large effect subjects (97.5%): ", round(exp(subject_slope[3]), 3), "x multiplier\n")
```

Large effect subjects (97.5%): 3.795 × multiplier


```
Visualization
hist(simulated_slopes,
 main = "Prior-implied Subject Random Slopes",
 xlab = "Condition effect adjustment (log-RT scale)",
 breaks = 30, col = "lightcoral", border = "white")
```



#### 1.4.3 Residual Noise Distribution

```
if (!is.null(sigma_vals) && is.numeric(sigma_vals)) {
 hist(sigma_vals,
 main = "Prior for Residual Noise (Sigma)",
 xlab = "Sigma value",
 breaks = 30, col = "lightgreen", border = "white")
} else {
 plot(1, main = "Sigma visualization", xlab = "", ylab = "")
 text(1, 1, "Sigma samples not available")
}
```

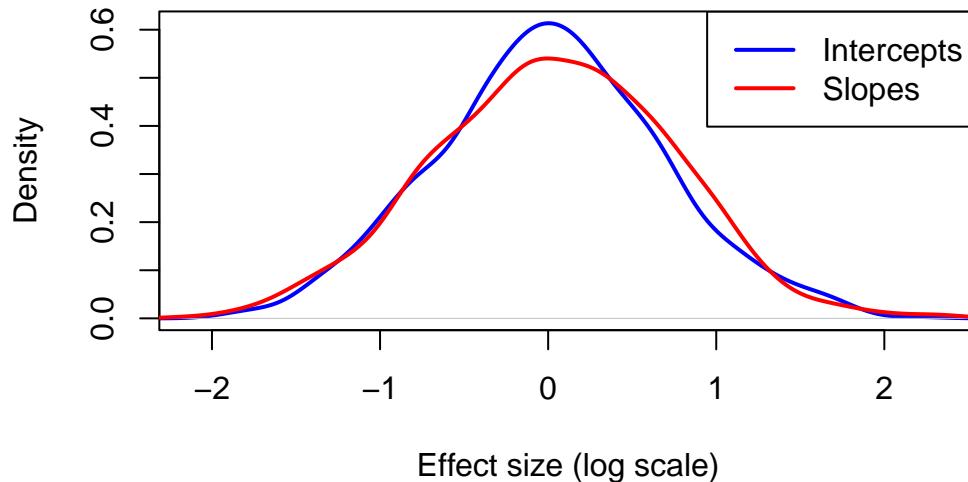
## Prior for Residual Noise (Sigma)



### 1.4.4 Random Effects Comparison

```
Density plot comparing intercepts and slopes
plot(density(simulated_intercepts),
 main = "Prior-implied Random Effects Distributions",
 xlab = "Effect size (log scale)",
 col = "blue", lwd = 2, xlim = range(c(simulated_intercepts, simulated_slopes)),
 ylim = c(0, max(density(simulated_intercepts)$y, density(simulated_slopes)$y)))
lines(density(simulated_slopes),
 col = "red", lwd = 2)
legend("topright", c("Intercepts", "Slopes"), col = c("blue", "red"), lwd = 2)
```

## Prior-implied Random Effects Distributions



## 1.5 Interpretation

### 1.5.1 Good Signs (Prior is Reasonable)

- Prior generates log-RTs around 6 (~400ms)
- 95% interval roughly 200-1100ms (plausible RT range)
- Condition effect typically < 150ms difference
- Between-subject variation is moderate

### 1.5.2 Problems to Watch For

- Mean RT  $\gg$  1000ms: intercept prior too high
- 95% interval 10ms-50s: priors too wide
- No variation between subjects: SD priors too small

## 1.6 Summary

Before fitting your model to actual data, always validate that your priors: 1. Generate reasonable predictions 2. Allow the data to inform the posterior 3. Respect domain knowledge constraints

Adjust priors as needed and rerun these checks.